

# EPOCH specific

“The Angry Penguin”, used under creative commons licence  
from Swantje Hess and Jannis Pohlmann.



Warwick RSE

# Why Fortran?

- Not dead language
  - Dominates in HPC sector
  - >70% of used core hours on ARCHER are Fortran
  - New standard being released this year
- Easier to develop in than C
- Faster execution than C++ (in general)

# Why Not Feature X?

- EPOCH has a very large user base
- Lots of people want lots of things
- Developer time is very limited
- Make feature requests on the forums
  - Remember that this is a *request*
  - It'll go on the list of things that we consider for future development

# Parallelism in EPOCH



# Domain Decomposition

- Split up the spatial domain so that each processor has it's own part
- Boundaries between processors are just like normal boundary conditions
- Just have to get the data from the neighbouring processor to populate the boundary
- Want to do as little communicating as possible (maximise computation done)
- Minimise surface area to volume ratio to get best scaling

# Domain Decomposition


- For fluid code (like CFSA's LARE code) job done
- Every processor has same sized part of domain
  - Finishes work at the same time
- For EPOCH have to worry about particles
  - Particles are free to move around
  - Can have more particles in one bit of grid than another

# Domain Decomposition

- Leads to a load balancing problem
- Have to adjust the amount of space on each processor so that they have the same number of particles, not the same amount of space
- Roughly anyway, things like the EM solver also takes work and that depends on number of grid cells
- How do you do load balancing?


# Load Balancing

- Only 1 processor has any work
- Have to load balance
- Simplest solution is to move processor edges "rigidly"
- Y range in domain can only depend on Y coordinate of processor
- And similarly for X






# Load Balancing




- Now all processors have same load
- Problem solved?
- Only for some cases
- This is EPOCH's current algorithm

# Load Balancing




- Not for this case!
- Two processors do almost no work at all
- Using the algorithm before load as balanced as possible
- How to fix this?

# Load Balancing



- Do it only in 1 dimension
- Can always split domain like this
- Performance drops as you have fewer grid cells on a processor
  - Communication dominates
- Limits maximum number of CPUs
- Is possible in current EPOCH

# Load Balancing



- Break the requirement that we move edges rigidly
- Load is now balanced much better for wider range of conditions
- Work starting to implement in EPOCH soon

IO in EPOCH

# SDF

- The SDF IO system is a low level self describing wrapper over the MPI-IO parallel output system
- File output is not dependent on number of processors that EPOCH was run on
- Performance is very high on well set up machines
- Gives us more control than using netCDF and a more seamless user experience than HDF5